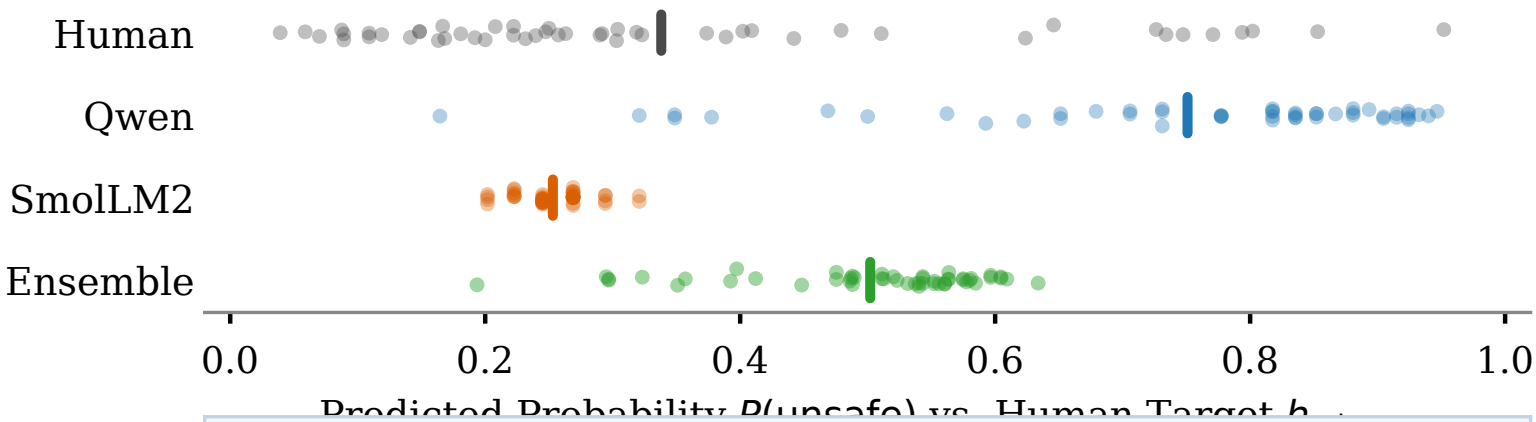

A. Qwen2.5-0.5B: Highest Judge Ranking Acuity

• AUROC: 0.7529, ρ: 0.5680 | Overprediction: $\mu = 0.751$ vs. 0.338 (MAE: 0.4152)

B. SmolLM2-1.7B: Lowest Judge Absolute Error

• Lowest Judge MAE: 0.1734 | Low Variance Anchor: $\mu = 0.253, \sigma = 0.028$ (AUROC: 0.6690)

C. Two-Judge Ensemble: Peak Rank Concordance & Trade-Offs

• Highest Rank Correlation: Spearman $\rho = 0.5903$, Pearson $r = 0.4992$

• MAE 0.2464 (-40.6% vs Qwen); Held-out baseline retains lower Soft MSE (0.0602)