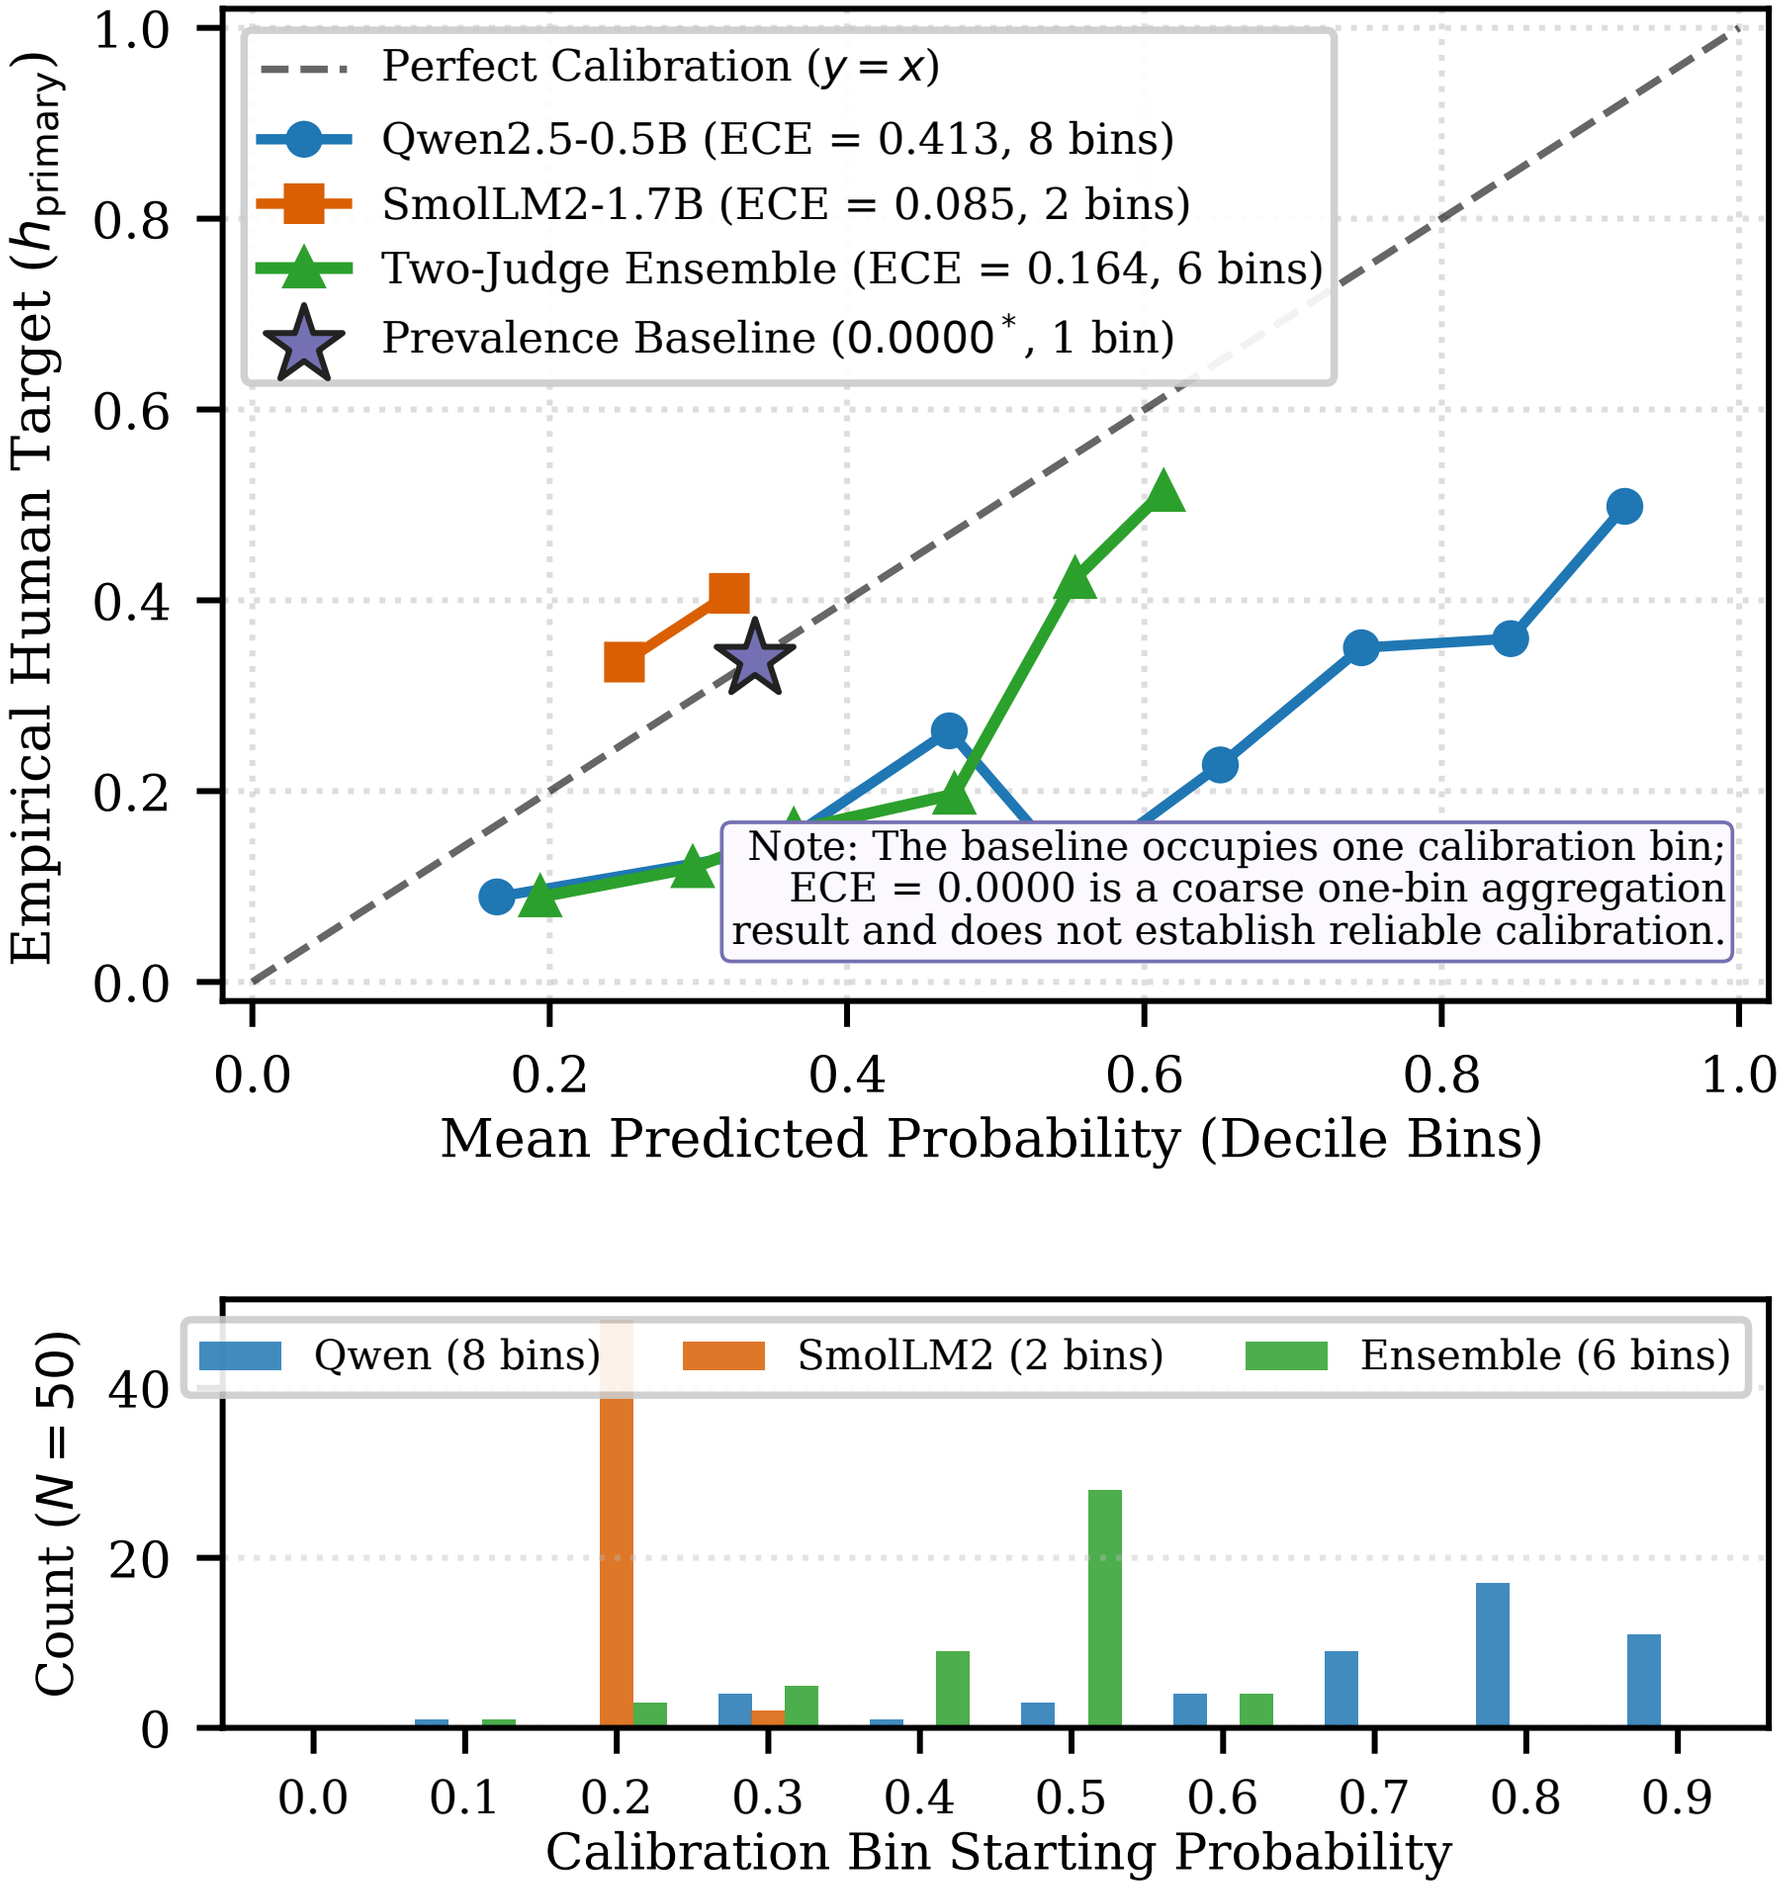
Perfect Calibration ($y = x$)
Qwen2.5-0.5B (ECE = 0.413, 8 bins)
SmolLM2-1.7B (ECE = 0.085, 2 bins)
Two-Judge Ensemble (ECE = 0.164, 6 bins)
Prevalence Baseline (0.0000*, 1 bin)
Note: The baseline occupies one calibration bin; ECE = 0.0000 is a coarse one-bin aggregation result and does not establish reliable calibration.
Empirical Human Target ($h_{\text{primary}}$)
Mean Predicted Probability (Decile Bins)
0.0
0.2
0.4
0.6
0.8
1.0
Qwen (8 bins)
SmolLM2 (2 bins)
Ensemble (6 bins)
Count ($N = 50$)
0
20
40
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
0.8
0.9
Calibration Bin Starting Probability